Distributed Algorithms (UAI/503): Berkeley Algorithm

Eldhose Poulose

January 7, 2023

Abstract

In this report, the Berkeley Algorithm and its implementation using Java RMI technology is explained. The source file can be viewed under Github.

1 Introduction

In the Distributed Systems (DS) the nodes are communicating with each other using message passing. To achieve real time applications working in ordered manner the reference parameter used is time. Therefore synchronisation of time is essential for allocating the available resources. Synchronization in DS can be achieved by using physical clock of the node. For synchronization purpose, each node in the system needs to share their local clock time with another node in the system.

2 Berkeley Algorithm

The Berkeley Algorithm follows a master slave communication operation. There is a central computer that serves as the master or a time server. The master master periodically sends a request message to all other slaves or nodes which asks the time of the destination nodes. The master will receive the round trip time value (RTT) from the slaves, and the master will average the time values including its own clock value and readjusts its own clock accordingly. The master will also eliminate readings from faulty clocks that is values far outside the range. It does this by taking a subset of the returned times with a small variance. After this, the master will then send the amount by which each individual's clock requires adjustment to each individual clock. This value can be positive or negative. If the master fails at any point, then a new master is elected to take over and function exactly like its predecessor.

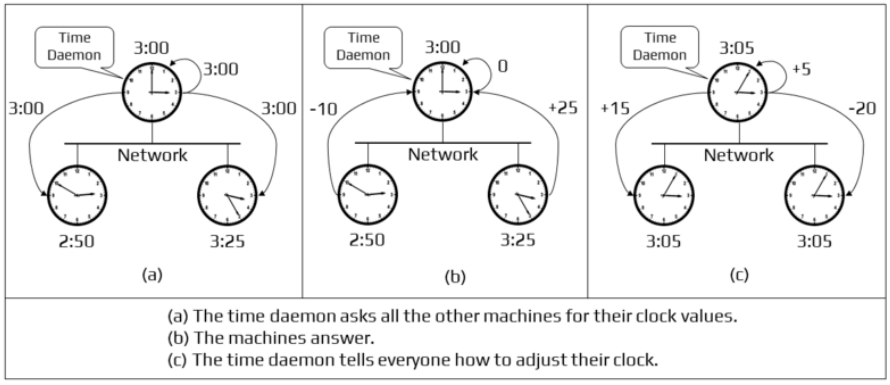

Figure 1: Berkeley Algorithm Workflow[source: see below]

3 Methods

- *client*: contains client side (master) implementation files
- *server*:

4 Results

To test the script run server.py on the machine that will receive the connections and client.py on the machine that will send the connections.

Libraries used in the script:

- time - time library
- socket - library for making connections to other hosts

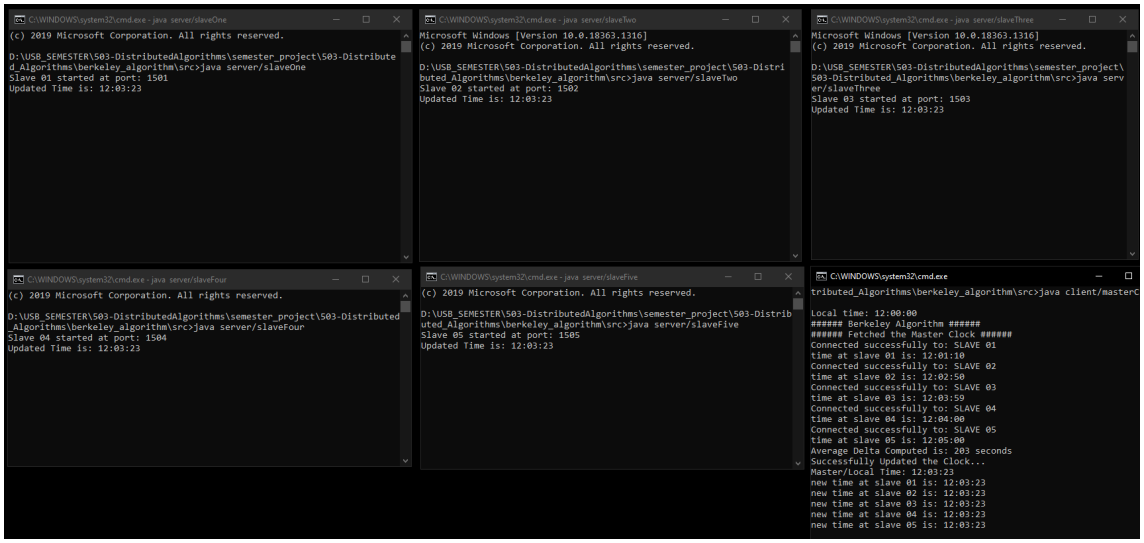

Figure 2: Updating 5 Slaves and a master clock time using Berkeley Algorithm

Link to Github

A source code: clock server

```
package server;

import java.rmi.Remote;
import java.rmi.RemoteException;
import java.time.LocalTime;


public interface clockServer extends Remote{

        LocalTime getTime() throws RemoteException;
        void adjustTime(LocalTime slaveTime, long delta) throws RemoteException;

}
```

B source code: clock server Implementation

```
package server;
import java.rmi.RemoteException;
import java.rmi.server.UnicastRemoteObject;
import java.time.LocalTime;
import global.berkeleyParams;

public class clockServerImplementation extends UnicastRemoteObject implements clockSe

        private LocalTime time;
        public clockServerImplementation(LocalTime time) throws RemoteException {
                this.time = time; // fetch the time
        }

        @Override
        public LocalTime getTime() throws RemoteException {
                return time;
        }

        @Override
        public void adjustTime(LocalTime slaveTime, long delta) throws RemoteExceptio
                long localTime = slaveTime.toSecondOfDay(); // fetch the time from sl
                long serverTime = this.getTime().toSecondOfDay(); // fetch the time fr
                var deltaTime = serverTime - localTime; // compute the difference in
                deltaTime = deltaTime * -1 + delta + serverTime;
                LocalTime updatedClock = LocalTime.ofSecondOfDay(deltaTime);
                System.out.println("Updated Time is: " + berkeleyParams.dateformatter
                this.time = updatedClock;
        }

}
```

C source code: global

```
//
package global;

import java.time.format.DateTimeFormatter;

public interface berkeleyParams {
```

```
        public final String servername01 = "localhost";
        public final int serverport01 = 1501;

        public final String servername02 = "localhost";
        public final int serverport02 = 1502;

        public final String servername03 = "localhost";
        public final int serverport03 = 1503;

        public final String servername04 = "localhost";
        public final int serverport04 = 1504;

        public final String servername05 = "localhost";
        public final int serverport05 = 1505;

        public final DateTimeFormatter dateformatter = DateTimeFormatter.ofPattern("HH

}
```

D source code: main

```
// This file includes the client side implementation of the algorithm
// 5 server/ slave machines time is fetched and the time is adjusted.
// time difference and the average of them is computed and adjusted


package client;
import java.rmi.*;
import java.rmi.registry.*;
import static global.berkeleyParams.dateformatter;
import java.rmi.registry.LocateRegistry;
import java.rmi.registry.Registry;
import java.time.LocalTime;
import java.util.ArrayList;
import global.berkeleyParams;
import server.clockServer;
import server.clockServerImplementation;

//master/Client Side Implementation

public class masterClock {

        public static void main(String[] args) {
                try {
                        var slavetimes = new ArrayList<LocalTime>();

                        LocalTime masterTime = LocalTime.parse("12:00:00", berkeleyPar
                        slavetimes.add(masterTime);
                        System.out.println("Local time: " + dateformatter.format(maste

                        System.out.println("###### Berkeley Algorithm ######");
                        System.out.println("###### Fetched the Master Clock ######");

                        // Connect to slave 01 and get Time
                        Registry registry01 = LocateRegistry.getRegistry(berkeleyParam
                        clockServer ts01 = (clockServer) registry01.lookup(clockServer
                        System.out.println("Connected successfully to: SLAVE 01");
                        LocalTime slaveTime01 = ts01.getTime();
```

```
slavetimes.add(slaveTime01);
System.out.println("time at slave 01 is: " + dateformatter.for

// Connect to slave 02 and get Time
Registry registry02 = LocateRegistry.getRegistry(berkeleyParan
clockServer ts02 = (clockServer) registry02.lookup(clockServer
System.out.println("Connected successfully to: SLAVE 02");
LocalTime slaveTime02 = ts02.getTime();
slavetimes.add(slaveTime02);
System.out.println("time at slave 02 is: " + dateformatter.for


// Connect to slave 03 and get Time
Registry registry03 = LocateRegistry.getRegistry(berkeleyParan
clockServer ts03 = (clockServer) registry03.lookup(clockServer
System.out.println("Connected successfully to: SLAVE 03");
LocalTime slaveTime03 = ts03.getTime();
slavetimes.add(slaveTime03);
System.out.println("time at slave 03 is: " + dateformatter.for

// Connect to slave 04 and get Time
Registry registry04 = LocateRegistry.getRegistry(berkeleyParan
clockServer ts04 = (clockServer) registry04.lookup(clockServer
System.out.println("Connected successfully to: SLAVE 04");
LocalTime slaveTime04 = ts04.getTime();
slavetimes.add(slaveTime04);
System.out.println("time at slave 04 is: " + dateformatter.for

// Connect to slave 04 and get Time
Registry registry05 = LocateRegistry.getRegistry(berkeleyParan
clockServer ts05 = (clockServer) registry05.lookup(clockServer
System.out.println("Connected successfully to: SLAVE 05");
LocalTime slaveTime05 = ts05.getTime();
slavetimes.add(slaveTime05);
System.out.println("time at slave 05 is: " + dateformatter.for

// COMPUTATION of TIME DIFFERENCE AND AVERAGING THEM

var localTimeSeconds = masterTime.toSecondOfDay();

var delta01 = slaveTime01.toSecondOfDay() - localTimeSeconds;
var delta02 = slaveTime02.toSecondOfDay() - localTimeSeconds;
var delta03 = slaveTime03.toSecondOfDay() - localTimeSeconds;
var delta04 = slaveTime04.toSecondOfDay() - localTimeSeconds;
var delta05 = slaveTime05.toSecondOfDay() - localTimeSeconds;

var delta_average = (delta01 + delta02 + delta03 + delta04 +
System.out.println("Average Delta Computed is: "+delta_average

// Assign new time
ts01.adjustTime(masterTime, delta_average);
ts02.adjustTime(masterTime, delta_average);
ts03.adjustTime(masterTime, delta_average);
ts04.adjustTime(masterTime, delta_average);
ts05.adjustTime(masterTime, delta_average);
masterTime = masterTime.plusSeconds(delta_average);
System.out.println("Successfully Updated the Clock...");
```

```
                        // Verifying the time on all machines
                        System.out.println("Master/Local Time: " + dateformatter.forma
                        System.out.println("new time at slave 01 is: " + dateformatter
                        System.out.println("new time at slave 02 is: " + dateformatter
                        System.out.println("new time at slave 03 is: " + dateformatter
                        System.out.println("new time at slave 04 is: " + dateformatter
                        System.out.println("new time at slave 05 is: " + dateformatter

                } catch (Exception ex) {
                        System.out.println(ex);
                }
        }

}
```

E source code: slave

```
package server;

import static global.berkeleyParams.dateformatter;
import java.rmi.registry.LocateRegistry;
import java.rmi.registry.Registry;
import java.time.LocalTime;
import global.berkeleyParams;

// Initiate the Slave 01

public class slaveOne {

        public static void main(String[] args) {
                try {
                        clockServer ts01 = new clockServerImplementation(LocalTime.pa
                        Registry registry01 = LocateRegistry.createRegistry(berkeleyPa
                        registry01.rebind(clockServerImplementation.class.getSimpleNam
                        System.out.println(String.format("Slave 01 started at port: %s
                } catch (Exception ex) {
                        System.out.println(ex);
                }
        }

}
```